
Abstract

We study a particular scheduling setting in which a set of n jobs with individual release times, deadlines, and processing volumes has to be scheduled across m homogeneous processors while minimizing the consumed energy. Idle processors can be turned off so as to save energy, while turning them on requires a fixed amount of energy. For the special case of a single processor, the greedy algorithm Left-to-Right guarantees an approximation factor of 2. We generalize this simple greedy policy to the case of multiple processors and show that the energy costs are still bounded by $2 \text{OPT} + P$, where P is the total amount processing volume. Our algorithm has a running time of $O(Fn \log d)$, where d is difference between the largest deadline and the earliest release time and F the running time required by a maximum-flow calculation for checking the feasibility of an instance.

Contents

1	Introduction	3
1.1	Busy and idle intervals	3
1.2	Lower and upper bounds for the amount of busy processors	3
1.3	Properties of an optimal schedule	3
2	Algorithm	5
3	Structure of the PLTR-Schedule	7
3.1	Preliminary Definitions	7
3.2	Critical set of time slots	7
3.3	Definitions based on critical sets	9
4	Modification of our Schedule	11
4.1	Augmentation	11
4.2	Realignment	11
4.3	Invariants for Realignment	12
4.4	Approximation Guarantee	16
5	Running Time	17

1 Introduction

We study a scheduling setting with a *power-down mechanism* and processors working in parallel. In this setting a system consisting of multiple homogeneous processors has to process a set of jobs while minimizing the energy consumed. Each job has an individual time interval during which it has to be processed. Idle processors can be turned off so as to save energy, while turning them on requires a fixed amount of energy. Intuitively one aims for few but long idle periods during which it is worthwhile to turn a processor off.

More specifically a problem instance consists of a set J of jobs with a release time r_j , deadline d_j , and processing volume p_j for every job $j \in J$. Each job $j \in J$ has to be scheduled for p_j units of time in the execution interval $E_j := [r_j, d_j]$ between its release time and its deadline, across $m \geq 1$ processors. Preemption of jobs and migration between processors is allowed at discrete times and occurs without delay, but no more than one processor may process any given job at the same time. Without loss of generality we assume the earliest release time r_{\min} to be 1 and denote the last deadline by d^* . The set of discrete time slots is denoted by $T := \{1, \dots, d^*\}$. The total amount of processing volume is $P := \sum_{j \in J} p_j$.

Every processor is either completely off or completely on in a discrete time slot $t \in T$. A processor can only work on some job in the time slot t if it is in the on-state. A processor can be turned on and off at discrete times without delay. All processors start in the off-state. The objective now is to find a feasible schedule which minimizes the expended energy E , which is defined as follows. Each processor consumes 1 unit of energy for every timeslot it is in the on-state and 0 units of energy if it is in the off-state. Turning a processor on consumes a constant amount of energy $q \geq 0$, which is fixed by the problem instance. In Graham's notation (Graham et al., 1979), this setting can be described with $m \mid r_j; \bar{d}_j; pmtn \mid E$.

1.1 Busy and idle intervals

We say a processor is *busy* at time t if some job is scheduled on this processor at time t . Otherwise, the processor is *idle*. Clearly a processor cannot be busy and off at the same time. Also note that if a processor is idle for more than q units of time, it is worth turning the processor off during this idle interval. Our algorithm will specify for each processor when it is busy and when it is idle. Each processor is then defined to be in the off-state during idle intervals of length greater than q and otherwise in the on-state.

1.2 Lower and upper bounds for the amount of busy processors

We specify a generalization of our problem which we call *power-down scheduling with lower and upper bounds*. Where in the original problem, for each time slot t , between 0 and m processors were allowed to be working on jobs, i.e. being busy, we now specify a lower bound $l_t \geq 0$ and an upper bound $m_t \leq m$. For a feasible solution to *power-down scheduling with lower and upper bounds*, we require that in every time slot t , the number of busy processors, which we will denote with $\text{vol}(t)$, lies within the lower and upper bounds, i.e. $l_t \leq \text{vol}(t) \leq m_t$. Note that this generalizes the problem of *deadline-scheduling-on-intervals* introduced by Antoniadis et al. (2020) by lower bounds.

1.3 Properties of an optimal schedule

Definition 1. *Given some arbitrary but fixed order on the number of processors, a schedule S fulfills the stair-property if it uses the lower numbered processors first, i.e. for every $t \in T$, if processor $k \in [m]$ is busy at t , then every processor $k' \leq k$ is busy at t . This symmetrically implies that if processor $k \in [m]$ is idle at t , then every processor $k' \geq k$ is idle at t .*

Lemma 2. *For every feasible problem instance, there exists an optimal schedule S_{opt} which fulfills the stair-property.*

Proof. Let S be an optimal schedule. We transform S such that it fulfills the stair-property without increasing its costs and while maintaining feasibility. Let $k, k' \in [m]$ be two processors with $k' < k$, job $j \in J$ scheduled on processor k in time slot $t \in T$ while k' is idle in t . Let I be the idle interval on processor k' containing t . We now move all jobs scheduled on processor k during I to be scheduled on processor k' instead. Since I is a maximal interval for which processor k' is idle, this modification does not increase the combined costs of processors k' and k . The modification also moves at least job

j from processor k down to k' while not moving any job from processor k' to k . Jobs are only moved between processor at the same time slot and only to slots of processor k' which are idle, hence the resulting schedule is still feasible. This modification can be repeated until the schedule has the desired property. \square

From now on, when considering an optimal schedule, we will silently assume that it fulfills the stair property based on Lemma 2.

Define busy, idle, off, on intervals.

2 Algorithm

The following algorithm *Parallel Left-to-Right* (PLTR) iterates through the processors in some arbitrary but fixed order and keeps the current processor idle for as long as possible such that the scheduling instance remains feasible. Once the current processor cannot be kept idle for any longer, it becomes busy and PLTR keeps it and all lower-numbered processors busy for as long as possible while again maintaining feasibility. The algorithm enforces these restrictions on the busy processors by iteratively making the upper and lower bounds m_t, l_t more restrictive. Visually, when considering the time slots on an axis from left to right and when stacking the schedules of the individual processors on top of each other, this generalization of the single processor *Left-to-Right* algorithm hence proceeds *Top-Left-to-Bottom-Right*.

Once the algorithm returns, an actual schedule can easily be constructed by running the flow-calculation used for the feasibility check depicted in Figure 2 or just taking the result of the last flow-calculation done during the algorithm. The mapping from the flow to an actual assignment of jobs to processors and time slots can then be defined as described in Lemma 3, which also ensures that the resulting schedule fulfills the stair-property from Definition 1, i.e. that it uses the lower-numbered processors first.

Algorithm 1 Parallel Left-to-Right

```

 $m_t \leftarrow m$  for all  $t \in T$ 
 $l_t \leftarrow 0$  for all  $t \in T$ 
for  $k \leftarrow m$  to 1 do
   $t \leftarrow 0$ 
  while  $t \leq d^*$  do
     $t \leftarrow \text{keepIdle}(k, t)$ 
     $t \leftarrow \text{keepBusy}(k, t)$ 
  
```

Algorithm 2 Auxiliary subroutines for PLTR

```

function keepIdle( $k, t$ )
  search for maximal  $t' > t$  s.t. exists feasible schedule with  $m_{t''} \leftarrow k - 1$  for all  $t'' \in [t, t')$ 
   $m_{t''} \leftarrow k - 1$  for all  $t'' \in [t, t')$ 
  return  $t'$ 

function keepBusy( $k, t$ )
  search for maximal  $t' > t$  s.t. exists feasible schedule with  $l_{t''} \leftarrow \max\{k, l_{t''}\}$  for all  $t'' \in [t, t')$ 
   $l_{t''} \leftarrow \max\{k, l_{t''}\}$  for all  $t'' \in [t, t')$ 
  return  $t'$ 
  
```

Formulate and format keepBusy and keepIdle nicer?

The feasibility check in subroutines keepIdle and keepBusy can be performed by calculating a maximum s - t flow in the flow network given in Figure 2.

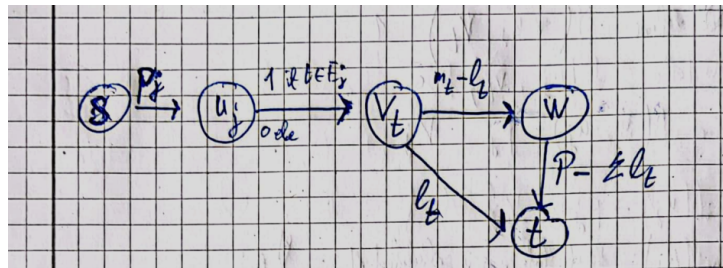


Figure 2.1: The Flow-Network for checking feasibility of a scheduling instance with lower and upper bounds l_t and m_t for the number of busy processors at t .

Lemma 3. *There exists a feasible solution to a scheduling instance with lower and upper bounds l_t, m_t if and only if the maximum s - t flow in the corresponding flow network depicted in Figure 2 has value P .*

Ambiguity of t used for time slots and sink in flow-network: Use α - ω flow

Proof. Let f be a s - t flow of value $|f| = P$. We construct a feasible schedule from f respecting the lower and upper bounds given by l_t and m_t . For every $j \in J$ and $t \in T$, if $f(u_j, v_t) = 1$, then schedule j at slot t on the lowest-numbered processor not scheduling some other job. Since $|f| = P$ and the capacity $c(\{s\}, V \setminus \{s\}) = P$, we have $f_{in}(u_j) = p_j$ for every $j \in J$. Hence $f_{out}(u_j) = \sum_{t \in E_j} f_{in}(v_t) = p_j$. Hence every job j is scheduled in p_j distinct time slots within its execution interval.

The schedule respects the upper bounds m_t , since $c(v_t, w) + c(v_t, t) \leq m_t - l_t + l_t$ and hence for every t at most m_t jobs are scheduled at t .

The schedule respects the lower bounds l_t , since $c(V \setminus \{t\}, \{t\}) = P$ and hence $f(v_t, t) = l_t$ for every slot t . By flow conservation we then have $f_{in}(v_t) \geq l_t$ which implies that at least l_t jobs are scheduled at every slot t .

For the other direction consider a feasible schedule respecting the lower and upper bounds l_t, m_t . We construct a flow f of value P and show that it is maximal.

If j is scheduled at slot t and hence $t \in E_j$, define $f(u_j, v_t) = 1$, otherwise $f(u_j, v_t) = 0$. Define $f(s, u_j) = p_j$ for every $j \in J$. Hence we have $f_{in}(u_j) = p_j$ and $f_{out}(u_j)$ must be p_j since this corresponds to the number of distinct time slots in which j is scheduled. Define $f(v_t, t) = l_t$ for every slot t . Define $f(v_t, w) = f_{in}(v_t) - l_t$. We have $f(v_t, w) \leq m_t - l_t$ since $f_{in}(v_t)$ corresponds to the number $\text{vol}(t)$ of jobs scheduled at t , which is at most m_t . We also have $f_{out}(v_t) = f_{in}(v_t) - l_t + l_t = f_{in}(v_t)$.

Define $f(w, t) = P - \sum_{t \in T} l_t$. Then $f_{in}(w) = \sum_{t \in T} f_{in}(v_t) - l_t = \sum_{t \in T} \text{vol}(t) - \sum_{t \in T} l_t$. Since the schedule is feasible, we have $\sum_{t \in T} \text{vol}(t) = P$ and finally the flow conservation $f_{in}(w) = P - \sum_{t \in T} l_t = f_{out}(w)$. □

Theorem 4. *Given a feasible problem instance, algorithm PLTR constructs a feasible schedule.*

Proof. By definition of subroutines keepIdle and keepBusy, PLTR only modifies the upper and lower bounds m_t, l_t for the number of busy processors such that the resulting instance of *power-down scheduling with lower and upper bounds* remains feasible. The correctness of the algorithm then follows from the correctness of the flow-calculation for checking feasibility, which is implied by Lemma 3. □

3 Structure of the PLTR-Schedule

3.1 Preliminary Definitions

Definition 5. For schedule S , we define the volume $\text{vol}_S(j, Q)$ of job $j \in J$ in a set $Q \subseteq T$ of time slots as the number of time slots of Q for which j is scheduled at by S .

Definition 6. We define the forced volume $\text{fv}(j, Q)$ of job $j \in J$ in a set $Q \subseteq T$ of time slots as the minimum number of time slots of Q for which j has to be scheduled in every feasible schedule, i.e.

$$\text{fv}(j, Q) := \max\{0; p_j - |E_j \setminus Q|\}.$$

Definition 7. We define the unnecessary volume $\text{uv}_S(j, Q)$ of job $j \in J$ in a set $Q \subseteq T$ of time slots as the amount of volume which does not have to be scheduled during Q , i.e.

$$\text{uv}_S(j, Q) := \text{vol}_S(j, Q) - \text{fv}(j, Q).$$

Definition 8. We define the possible volume $\text{pv}(j, Q)$ of job $j \in J$ in a set $Q \subseteq T$ of time slots as the maximum amount of volume which j can be feasibly scheduled in Q , i.e.

$$\text{pv}(j, Q) := \min\{p_j, |E_j \cap Q|\}.$$

Since the corresponding schedule S will always be clear from context, we drop the subscript for vol and uv . We extend our volume definitions to single timeslots $t \in T$ and to sets $J' \subseteq J$ of jobs by summing over all $j \in J'$, i.e.

$$\begin{aligned} \text{vol}(J', Q) &:= \sum_{j \in J'} \text{vol}(j, Q), \\ \text{vol}(t) &:= \text{vol}(J, \{t\}). \end{aligned}$$

If the first parameter is dropped, we refer to the whole set J , i.e. $\text{vol}(Q) := \text{vol}(J, Q)$. Clearly we have for every feasible schedule, every $Q \subseteq T, j \in J$ that $\text{fv}(j, Q) \leq \text{vol}(j, Q) \leq \text{pv}(j, Q)$.

Definition 9. We define the density $\phi(Q)$ for a set $Q \subseteq T$ as the average amount of processing volume which has to be completed in every slot of Q , i.e. $\phi(Q) := \text{fv}(J, Q)/|Q|$. We also define $\hat{\phi}(Q) := \max_{Q' \subseteq Q} \phi(Q')$.

If $\hat{\phi}(Q) > k - 1$, then clearly at least k processors are required in some time slot $t \in Q$ for every feasible schedule.

Definition 10. We define the deficiency $\text{def}(Q)$ of a set $Q \subseteq T$ of time slots as the difference between the amount of volume which has to be completed in Q and the processing capacity available in Q , i.e. $\text{def}(Q) := \text{fv}(Q) - \sum_{t \in Q} m_t$.

Definition 11. We define the excess $\text{exc}(Q)$ of a set $Q \subseteq T$ of time slots as the difference between the processor utilization required in Q and the amount of processing volume available in Q , i.e. $\text{exc}(Q) := \sum_{t \in Q} l_t - \text{pv}(Q)$.

3.2 Critical set of time slots

Lemma 12. For every s - t cut (S, \bar{S}) in the network given in Figure 2 we have at least one of the following two lower bounds for the capacity $c(S)$ of the cut: $c(S) \geq P - \text{def}(Q(S))$ or $c(S) \geq P - \text{exc}(Q(\bar{S}))$, where $Q(S) := \{t \mid v_t \in S\}$.

Proof. Let (S, \bar{S}) be an s - t cut, let $J(S) := \{j \mid u_j \in S\}$. We consider the contribution of every node of S to the capacity $c(S)$ of the cut. First consider the case that $w \notin S$.

- Node s : $\sum_{j \in J(\bar{S})} p_j$.
- Node u_j : $|\{v_t \in \bar{S} \mid t \in E_j\}| = |E_j \setminus Q(S)| \geq p_j - \text{fv}(j, Q(S))$
- Node v_t : $l_t + m_t - l_t = m_t$

Use just one definition environment for all volume definitions?

The inequality for node u_j follows since $\text{fv}(j, Q(S)) = \max\{0, p_j - |E_j \setminus Q(S)|\}$. In total, we can bound the capacity from below with

$$\begin{aligned} c(S) &\geq \sum_{j \in J(\bar{S})} p_j + \sum_{j \in J(S)} p_j - \text{fv}(j, Q(S)) + \sum_{t \in Q(S)} m_t \\ &= P - \text{fv}(J(S), Q(S)) + \sum_{t \in Q(S)} m_t \\ &\geq P - \text{def}(Q(S)). \end{aligned}$$

If $w \in S$, we have the following contributions of nodes in S .

- Node s : $\sum_{j \in J(\bar{S})} p_j \geq \text{pv}(J(\bar{S}), Q(\bar{S}))$.
- Node u_j : $|E_j \setminus Q(S)| = |E_j \cap Q(\bar{S})| \geq \text{pv}(j, Q(\bar{S}))$
- Node v_t : l_t
- Node w : $P - \sum_{t \in T} l_t$

In total, we have the lower bound

$$\begin{aligned} c(S) &\geq P + \text{pv}(Q(\bar{S})) - \sum_{t \in Q(\bar{S})} l_t \\ &= P - \text{exc}(Q(\bar{S})). \end{aligned}$$

□

Lemma 13. *An instance of power-down scheduling with lower and upper bounds is feasible if and only if $\text{def}(Q) \leq 0$ and $\text{exc}(Q) \leq 0$ for every $Q \subseteq T$.*

Proof. If $\text{def}(Q) > 0$ for some Q , then some upper bound m_t cannot be met. If $\text{exc}(Q) > 0$ for some Q , then some lower bound l_t cannot be met. For the direction from right to left, consider an infeasible scheduling instance with lower and upper bounds. By Lemma 3 we have that the maximum flow f for this instance has value $|f| < P$. Hence, there must be a s - t cut (S, \bar{S}) of capacity $c(S) < P$. Lemma 12 now implies that $\text{def}(Q(S)) > 0$ or $\text{exc}(Q(\bar{S})) > 0$. □

Definition 14. *A time slot $t \in T$ is called activation of processor $k \in [m]$ if processor k is idle at $t - 1$ and busy at t . A time slot $t \in T$ is just called activation if it is an activation of processor k for some $k \in [m]$.*

The following lemma provides the crucial structure required for the proof of the approximation guarantee. Intuitively, it states that for every activation t of processor k in the schedule S_{pltr} returned by PLTR, there must be some set Q of timeslots during which an activation of processor k is necessary in every feasible schedule.

Lemma 15. *For every time slot $t \in T$ for which some processor $k \in [m]$ becomes busy in S_{pltr} , there exists a set $Q \subseteq T$ of time slots with $t \in Q$,*

$$\begin{aligned} \text{fv}(Q) &= \text{vol}(Q), \\ \text{vol}(t') &\geq k - 1 && \text{for all } t' \in Q \text{ and} \\ \text{vol}(t') &\geq k && \text{for all } t' \in Q \text{ with } t' \geq t. \end{aligned}$$

Proof. Suppose for contradiction there is some activation $t \in T$ of processor $k \in [m]$ and no such Q exists for t . We show that PLTR would have extended the idle interval on processor k which ends at t . Consider the step in PLTR when t was the result of `keepIdle` on processor k and the corresponding lower and upper bounds $m_{t'}, l_{t'}$ for $t' \in T$ right after the calculation of t with the corresponding update of the bounds by `keepIdle`. We modify the bounds by decreasing m_t by 1. Note that at this point $m_{t'} \geq k$ for every $t' > t$ and $m_{t'} \geq k - 1$ for every t' .

Consider $Q \subseteq T$ s.t. $t \in Q$ and $\text{fv}(Q) < \text{vol}(Q)$. Before our decrement of m_t we had $m_Q := \sum_{t' \in Q} m_{t'} \geq \text{vol}(Q) > \text{fv}(Q)$. The inequality $m_Q \geq \text{vol}(Q)$ here follows since the upper bounds $m_{t'}$ are

How to replace the term activation? Candidates: start-of-work.

monotonically decreasing during PLTR. Since our modification decreases m_Q by at most 1, we hence still have $m_Q \geq \text{fv}(Q)$ after the decrement of m_t .

Consider $Q \subseteq T$ s.t. $t \in Q$ and $\text{vol}(t') < k - 1$ for some t' . At the step in PLTR considered by us, we hence have $m_{t'} \geq k - 1 > \text{vol}(t')$. Before our decrement of m_t we therefore have $m_Q > \text{vol}(Q) \geq \text{fv}(Q)$, which implies $m_Q \geq \text{fv}(Q)$ after the decrement.

Finally, consider $Q \subseteq T$ s.t. $t \in Q$ and $\text{vol}(t') < k$ for some $t' > t$. At the step in PLTR considered by us, we again have $m_{t'} \geq k > \text{vol}(t')$, which implies $m_Q \geq \text{fv}(Q)$ after our decrement of m_t . In summary, if for t no Q exists as characterized in the proposition, the activation of processor k at t could not have been the result of keepIdle on processor k . \square

Definition 16. We call such $Q \subseteq T$ for activations t of processor k characterized by Lemma 15 tight set Q_t over activation t of processor k .

Definition 17. A critical set $C_t \subseteq T$ over an activation t is the maximum of the set of tight sets Q_t over activation t in regard to the density ϕ , i.e.

$$C_t := \text{argmax}\{\phi(Q) \mid Q \subseteq T \text{ is tight set over } t\}.$$

As the set of these critical sets C_t for fixed t is closed under union, we take C_t to be the inclusion-maximal critical set over activation t for the sake of uniqueness.

3.3 Definitions based on critical sets

Definition 18. We define a total order \succeq on the set of critical sets C_t across all activations t . For activations $t, t' \in T$ of processors k and k' respectively, we define $C_t \succeq C_{t'}$ if and only if $k > k'$ or $k = k'$ and $t \leq t'$. In other words, \succeq defines the same order in which PLTR calculates the activations: from Top-Left to Bottom-Right. Equality in regard to \succeq is denoted with \sim .

Definition 19. Let $\text{rank} : \{C_t\} \rightarrow \mathbb{N}$ be a mapping to the natural numbers corresponding to \succeq , i.e.

$$\text{rank}(C_t) \geq \text{rank}(C_{t'}) \Leftrightarrow C_t \succeq C_{t'}.$$

Definition 20. Let $\text{crit} : \{C_t\} \rightarrow [m]$ be a mapping to the processors s.t.

$$\text{crit}(C_t) = c \Leftrightarrow c \text{ is the highest processor activated at } t.$$

Definition 21. We extend these definitions to general time slots $t \in T$ as follows.

$$\text{rank}(t) := \begin{cases} \max\{\text{rank}(C) \mid t \in C\} & \text{if } t \in C \text{ for some critical set } C \\ 0 & \text{otherwise} \end{cases}$$

$$\text{crit}(t) := \begin{cases} \max\{\text{crit}(C) \mid t \in C\} & \text{if } t \in C \text{ for some critical set } C \\ 0 & \text{otherwise} \end{cases}$$

We also extend the definitions to intervals $D \subseteq T$ with

$$\begin{aligned} \text{rank}(D) &:= \max\{\text{rank}(t) \mid t \in D\}, \\ \text{crit}(D) &:= \max\{\text{crit}(t) \mid t \in D\}. \end{aligned}$$

Definition 22. Let C be a critical set. A nonempty interval $V \subseteq T$ is a valley of $\text{rank}(C)$ if $C \sim V$ and V is inclusion maximal. Let C_1, \dots, C_l be the intervals contained in C . A nonempty interval V is a valley of C if V is exactly the interval between C_a and C_{a+1} for some $a < l$, i.e. $V = [\max C_a + 1, \min C_{a+1} - 1]$.

Definition 23. For a critical set C , an interval D spans C if $D \cap C$ contains only full subintervals of C and at least one subinterval of C . The left valley V_l of C and an interval D spanning C is the valley of C ending at $\min(C \cap D) - 1$ if such a valley of C exists. Symmetrically, the right valley V_r of C and an interval D spanning C is the valley of C starting at $\max(C \cap D) + 1$ if such a valley of C exists.

Definition 24. For a valley V , we define the jobs $J_V \subseteq J$ as all jobs which are scheduled by S_{pltr} in every $t \in V$.

Think about using a different term than to span here, e.g. to partially span.

Use $J(V)$ everywhere to not have double subscripts later?

Lemma 25. *For every critical set C with $c := \text{crit}(C)$, every interval $D \subseteq T$ spanning C : if $\phi(C \cap D) \leq c - \delta$ for some $\delta \in \mathbb{N}$, then the left valley V_l or the right valley V_r of C and D is defined and $|J_{V_l}| + |J_{V_r}| \geq \delta$, where we take $|J_V| := 0$ if V is not defined.*

Proof. By choice of C as critical set with $c = \text{crit}(C)$ we have $\text{vol}(C \cap D) \geq (c - 1) \cdot |C \cap D|$. If this inequality is fulfilled strictly, i.e. if $\text{vol}(C \cap D) > (c - 1) \cdot |C \cap D|$, then with the premise $\text{fv}(C \cap D)/|C \cap D| \leq c - \delta$ we directly get $\text{uv}(C \cap D)/|C \cap D| > \delta - 1$. This implies that there are at least δ jobs j scheduled in $C \cap D$ with $\text{uv}(j, C \cap D) > 0$. Such jobs can be scheduled in the part of C not contained in D , i.e. we must have $E_j \cap (C \setminus D) \neq \emptyset$ and hence the left valley V_l or the right valley V_r of C and D must exist. These jobs must be contained in J_{V_l} or J_{V_r} since they are scheduled in C and the schedule is feasible.

If on the other hand we have equality, i.e. $\text{vol}(C \cap D) = (c - 1) \cdot |C \cap D|$, then let t be the activation of processor c for which C is the critical set. Since $\text{vol}(t) > c - 1$, we must have $t \notin C \cap D$. By the same argument as before, we have that if $\text{fv}(C \cap D)/|C \cap D| \leq c - \delta$, then $\text{uv}(C \cap D)/|C \cap D| \geq \delta - 1$. Let $J' := \{j \in J \mid \text{uv}(j, C \cap D) > 0\}$. Clearly $|J'| \geq \delta - 1$. If this lower bound is fulfilled with equality, every $j \in J'$ must be scheduled in every time slot of $C \cap D$. Now suppose for contradiction that all jobs j scheduled during C which are not contained in J' have $E_j \cap C \cap D = \emptyset$. Then $\text{fv}(C \setminus D) = \text{vol}(C \setminus D)$ and we get $\phi(C \setminus D) > \phi(C)$ since by case assumption $\text{vol}(C \cap D) = (c - 1) \cdot |C \cap D|$. In conclusion, $C \setminus D$ is still a tight set over t but has higher density than C , contradicting the choice of C . Therefore, there must exist a job $j \notin J'$ scheduled in C which can be scheduled $C \cap D$.

In total, we have at least δ jobs scheduled in C with an execution interval intersecting both $C \setminus D$ and $C \cap D$. This implies that the left valley V_l or the right valley V_r of C and D exists and that at least δ jobs are contained in J_{V_l} or J_{V_r} . \square

Provide a rough visual sketch here

Make math more readable in this whole proof, e.g. by using fractions and display math, by replacing division by multiplication, or by introducing $t := |C \cap D|$

4 Modification of our Schedule

We modify the schedule S_{pltr} returned by PLTR in two steps. The first step augments specific processors with auxiliary busy slots s.t. in every critical set C at least the first $\text{crit}(C)$ processors are busy all the time. For the single processor Left-to-Right algorithm, the crucial property for the approximation guarantee is that every idle interval of S_{opt} can intersect at most 2 distinct idle intervals of the schedule returned by Left-to-Right. The second modification step of S_{pltr} is more involved and establishes this crucial property on every processor $k \in [m]$ by making use of Lemma 25. More specifically, it will establish that $\hat{\phi}(A) > k - 1$ for every busy interval A on processor k , i.e. that every feasible schedule requires k busy processors at some point during A .

These modification steps are only done for the sake of the analysis of the original schedule S_{pltr} and are not part of the algorithm. By making sure that the modifications cannot decrease the costs of our schedule, we get an upper bound for the costs of S_{pltr} .

4.1 Augmentation

We transform S_{pltr} into the *augmented schedule* S_{aug} by adding for every t with $k := \text{crit}(t) \geq 2$ and $\text{vol}(t) = k - 1$ an auxiliary busy slot on processor k . No job is scheduled in this auxiliary busy slot on processor k and it does also not count towards the volume of this slot. It merely forces processor k to be in the on-state at time k and allows us to keep thinking in terms of idle and busy intervals in our analysis of the costs.

Make this into a proper definition of S_{aug} ?

Lemma 26. *In S_{aug} processors $1, \dots, \text{crit}(t)$ are busy in every slot $t \in T$ with $\text{crit}(t) \geq 2$.*

Proof. The property directly follows from our choice of the critical sets, the definition of $\text{crit}(t)$ and the construction of S_{aug} . \square

4.2 Realignment

The intuition of the realignment of S_{aug} is the following. Lemma 26 guarantees us that every busy interval A on processor k spans the critical set C with $C \sim A$. It also guarantees that the left and right valley V_l, V_r of C and A do not end within an idle interval on processor k , provided that they are defined. Lemma 25 in turn guarantees us that if the density of A does not guarantee that S_{opt} has to use processor k during A , i.e. if $\hat{\phi}(A) \leq k - 1$, then V_l or V_r is defined and there is some j scheduled in every slot of V_l or V_r . Let V be the corresponding left or right valley of C and D for which such a job exists. Instead of scheduling this job on the processors below k , we can schedule the job on processor k in idle time slots during V instead of the lower processors, merging the busy interval A with at least one neighbouring busy interval on processor k . In the definition of the realignment, we will call this process of filling the idle slots during V on processor k closing of valley V on processor k .

The crucial part is making sure that this realignment continues to be possible whenever we have a busy interval with a density which is too small. For this purpose, we go through the busy intervals on each processor in the order of their rank, i.e. in the order of \succeq . We also allow every job to be used twice for the realignment by introducing further auxiliary busy slots (for an interval D spanning the critical set C , both the right and the left valley might be closed in the worst case). This allows us to maintain the Invariants stated in Lemma 27 during the realignment process, which are analogous to Lemma 25 and Lemma 26.

Algorithm 3 Realignment of S_{aug}

```

Res( $V$ )  $\leftarrow 2|J_V|$  for every valley  $V$ 
for  $k \leftarrow m$  to 1 do
    fill( $k, T$ )
    Res( $V$ )  $\leftarrow$  Res( $V$ )  $- 1$  for every valley  $V$  s.t. some  $V'$  with  $V' \cap V \neq \emptyset$  was closed on processor  $k$ 

```

Explicitly define valley V having been closed on processor k as $\text{close}(k, V)$ having been called.

Algorithm 4 Subroutines for the Realignment of S_{aug}

```

function fill( $k, V$ )
  if crit( $V$ )  $\leq 1$  then
    return
  let  $C$  be the critical set s.t.  $C \sim V$ 
  while exists busy interval  $A \subseteq V$  on processor  $k$  with  $A \sim V$  and  $\hat{\phi}(A) \leq k - 1$  do
    let  $V_l, V_r$  be the left and right valley for  $C$  and interval  $A$  (if  $A$  spans  $C$ )
    if  $V_l$  exists and Res( $V_l$ )  $> 0$  then
      close( $k, V_l$ )
    else if  $V_r$  exists and Res( $V_r$ )  $> 0$  then
      close( $k, V_r$ )
  for every valley  $V' \subseteq V$  of  $C$  which has not been closed on  $k$  do
    fill( $k, V'$ )
function close( $k, V$ )
  for every  $t \in V$  which is idle on processor  $k$  do
    if processors  $1, \dots, k - 1$  idle at  $t$  then
      introduce new auxiliary busy slot on processor  $k$  at time  $t$ 
    else
      move busy slot at time  $t$  of highest processor among  $1, \dots, k - 1$  to processor  $k$  at  $t$ 

```

4.3 Invariants for Realignment

Lemma 27. *For an arbitrary step during the realignment of S_{aug} and a valley $V \subseteq T$, let the critical processor k_V for V be the highest processor s.t.*

- *processor k_V is not fully filled yet, i.e. fill(k_V, T) has not yet returned,*
- *no $V' \supseteq V$ has been closed on k_V so far and*
- *there is a (full) busy interval $A \subseteq V$ on processor k_V .*

We take $k_V := 0$ if no such processor exists. At every step in the realignment of S_{aug} the following invariants holds for every valley V , where C denotes the critical set with $C \sim V$.

1. *If $\phi(C \cap D) \leq k_V - \delta$ for some $\delta \in \mathbb{N}$, some interval $D \subseteq V$ spanning C , then the left valley V_l or the right valley V_r of C, D exists and $\text{Res}(V_l) + \text{Res}(V_r) \geq 2\delta$.*
2. *For every $t \in C \cap V$, processors $1, \dots, k_V$ are busy at t .*
3. *Every busy interval $A \subseteq V$ on processor k_V with $A \sim V$ spans C .*

Proof. We show Invariants 1 and 2 via structural induction on the realigned schedule S_{real} . Then we show that Invariant 2 implies Invariant 3. For the induction base, consider S_{aug} , let V be an arbitrary valley in S_{aug} with $\text{crit}(V) \geq 2$ and let C the critical set with $C \sim V, c := \text{crit}(V)$.

We must have $k_V \leq c$, otherwise V would contain a full busy interval on processor $k_V > c$ and hence also an activation $t \in V$ of processor k_V , which by construction of S_{aug} would have $\text{crit}(t) = k_V > c$. This is a direct contradiction to $\text{crit}(V) = \max_{t \in V} \text{crit}(t) = c$.

Invariant 2 now follows since by construction of S_{aug} and our choice of C we have for every $t \in C$ that processors $1, \dots, k_V, \dots, c$ are busy at t . For Invariant 1, let D be an interval spanning C with $\phi(C \cap D) \leq k_V - \delta$ for some $\delta \in \mathbb{N}$. With $k_V \leq c$ we get $\phi(C \cap D) \leq c - \delta$ and hence by Lemma 25, we have that the left valley V_l or the right valley V_r of C and D exists and $|J_{V_l}| + |J_{V_r}| \geq \delta$. With the initial definition of Res(V) we get the desired lower bound of $\text{Res}(V_l) + \text{Res}(V_r) \geq 2\delta$.

Now suppose that Invariants 1 and 2 hold at all steps of the realignment up to a specific next step. Let V again be an arbitrary valley of $\text{crit}(V) \geq 2$ and k the processor currently being filled. Let furthermore k_V, k'_V be the critical processor for V before and after, respectively, the next step in the realignment. There are four cases to consider for this next step.

Case 1: Some $V' \supseteq V$ is closed on processor k . Then no valley W intersecting V has been closed so far on k . Also, since $\text{close}(k, \cdot)$ only moves the busy slot of the highest busy processor below k , we know that the stair property holds within V on processors $1, \dots, k$. We show that the closing of V' on k reduces the critical processor of V by at least 1, i.e. $k'_V \leq k_V - 1$. If $k_V = k$, then $V' \supseteq V$ is closed on processor k_V and hence by definition we have $k'_V \leq k_V - 1$. If $k_V < k$, suppose for contradiction that $k_V \leq k'_V \leq k$, where $k'_V \leq k$ again by definition since $V' \subseteq V$ is closed on processor k .

define stair property

Let $A \subseteq V$ be a full busy interval on k_V before the close of V' . We show that $A \subset V$, i.e. that there must be some $t \in V$ idle on k_V before the close and hence by the stair-property processors k_V, \dots, k idle at t before the close. If $V' \supseteq V$ is closed, clearly $V \subset T$ by the choice of V_l and V_r as valleys of some critical set in the realignment definition. Hence we know that $\min V - 1 \in T$ and $\max V + 1 \in T$. We show that $t := \min V - 1$ must be busy on k_V before the close. Let $W \supseteq V$ be the valley with $W \sim t$ and $t \in W$. We know that $W \supseteq V$ since $W \sim t \succ V$ since V is a valley. By our case assumption, no $W' \subseteq W$ can have been closed on processor k so far. With $W \supseteq V$ and the definition of k_W we get $k_W \geq k_V$, where k_W is the critical processor of W before the close. Our induction hypothesis now implies that processors $1, \dots, k_V, \dots, k_W$ are busy at t before the close. For $A \subseteq V$ to be a (full) busy interval on k_V before the close, we hence must have $\min V \notin A$. We know by definition of the realignment and function close, that for every k' with $k_V \leq k' < k$ and every $t \in V$:

Double check if this is true (or if we only have one of the two guaranteed): Since V must be valley of some C , we should have both $0 \notin V$ and $d^* \notin V$.

- If t was idle on k' before the close, then t is still idle on k' after the close (definition of close, $k' < k$).
- If t was idle on k_V before the close, then t was idle on k' before (stair-property) and hence t is still idle on k' after the close.
- If t was part of full a busy interval $A \subset V$ on k_V before the close, then t was idle on $k_V + 1$ before the close (choice of k_V). Hence t was idle on k before (stair-property) and hence t is idle on k_V after the close.

Taken together, for $t \in V$ to be busy on k' after the close, t must have been busy on k' before the close (definition close, $k' < k$) and t cannot have been part of a full busy interval $A \subseteq V$. Hence $t \in A$ for some partial busy interval $A \subseteq V$ on k' before the close. For $A' \subseteq V$ to be a full busy interval on k'_V after the close (with $k_V \leq k'_V < k$), we must have $A' \subseteq A$.

properly define partial busy intervals




Figure 4.1: The situation for case 1 in Lemma 27.

Hence there must have been a busy interval $A'' \subseteq [\min A', \max A]$ on processor $k'_V + 1 > k_V$ before the close, which contradicts the definition of $k_V < k$. In conclusion, we have $k'_V \leq k_V - 1$ which allows us to prove our invariants 1 and 2. If $\phi(C \cap D) \leq k'_V - \delta$ for some $\delta \in \mathbb{N}$ and some interval D spanning C , then $\phi(C \cap D) \leq k_V - (\delta + 1)$ and hence by induction hypothesis the left valley V_l or the right valley V_r for C, D exists and $\text{Res}(V_l) + \text{Res}(V_r) \leq 2(\delta + 1)$ both before and after the close. Our induction hypothesis also implies that for every $t \in C \cap V$ processors $1, \dots, k_V$ are busy before the close. Since at most the uppermost busy slot is moved by close, after the close of V' we still have processors $1, \dots, k_V - 1 \geq k'_V$ busy.

Case 2: Some $V' \subset V$ is closed on processor k . Again, no $V'' \supseteq V$ can have been close on processor k so far. We show that $k_V = k \geq k'_V$, i.e. that the critical processor of V before the close of V' is the processor currently being filled. Let W be the valley for which V' is closed, i.e. V' is closed during $\text{fill}(k, W)$. We must have $W \supset V'$ and therefore no $W' \supseteq W$ has been closed on k so far. Also, for V' to be closed in $\text{fill}(k, W)$, there must be some busy interval $A \subseteq W$ on k before the close, hence $k_W = k$. Since $V' \subset V$ and $V' \subset W$, V and W intersect ($V' \neq \emptyset$ by definition of V' as valley). If $V \prec W$, then by

choice of V' as valley of C_W we must have $V \subseteq V'$, which contradicts our case assumption. Therefore $V \succeq W$ and $V \supseteq W$ implying $k_V \leq k_W = k$ and since processor $k+1$ is already filled before the close we have $k_V = k \geq k'_V$.

For invariant 1, again let $\phi(C \cap D) \leq k'_V - \delta \leq k_V - \delta$ for some $\delta \in \mathbb{N}$ and some interval D spanning C . Our induction hypothesis implies that the left valley V_l or the right valley V_r of C, D exists and that both before and after the close we have $\text{Res}(V_l) + \text{Res}(V_r) \geq 2\delta$.

For invariant 2, observe that $V' \cap C = \emptyset$ since $V' \prec C$ by our case assumption. Our invariant 2 now directly follows from the induction hypothesis and $k'_V \leq k_V$.

Case 3: Some V' with $V' \cap V = \emptyset$ is closed on processor k . We first show that $\min V - 1 \notin V'$ and symmetrically $\max V + 1 \notin V'$. Consider $t := \min V - 1$. We know $t \in T$ since V' is valley of some critical set. By choice of V and t we must have $t \succ V$. If $t \in V'$, we would have $V' \succ V$ and hence $V' \supseteq V$, which contradicts our case assumption. Symmetrically, we know that $\max V + 1 \notin V'$. Therefore the close of V' does not modify the schedule within $[\min V - 1, \max V + 1]$, implying that no partial busy interval in V before the close can become a full busy interval. In summary we have $k_V = k'_V$ and invariants 1 and 2 follow as in case 2.

Case 4: The call to $\text{fill}(k, T)$ returns and $\text{Res}(V')$ is decreased by 1 for every V' such that some valley intersecting V' has been closed during $\text{fill}(k, T)$. First observe that the schedule itself does not change by this step but processor k is now fully filled, which implies $k'_V \leq k_V$. Invariant 2 then follows directly from the induction hypothesis.

We consider two subcases. If during $\text{fill}(k, T)$, no valley V' intersecting V was closed on k , then $\text{Res}(V)$ does not change and Invariant 1 follows from the induction hypothesis and $k'_V \leq k_V$. If on the other hand some valley V' intersecting V was closed on k during $\text{fill}(k, T)$, then $\text{Res}(V)$ is decreased by 1 to $\text{Res}'(V) := \text{Res}(V) - 1$. As seen in cases 1 to 3, k_V decreases monotonically during $\text{fill}(k, T)$. Consider the schedule right before the first valley V' intersecting V is closed on k . Let k_V^0 be the critical processor for V at this point of the realignment and k_V^1 the critical processor right after V' is closed.

We have $k'_V \leq k_V^0 - 1$: If $V' \supseteq V$, then as argued in case 1 $k_V^1 \leq k_V^0 - 1$ and hence $k'_V \leq k_V \leq k_V^1 \leq k_V^0 - 1$. If $V' \subset V$, then as argued in case 2 we have $k_V^0 = k$. Since $\text{fill}(k, T)$ returns in the next step of our case assumption, we have $k'_V \leq k - 1$ and hence $k'_V \leq k_V^0 - 1$. Invariant 2 now follows by our (strong) induction hypothesis. If $\phi(C \cap D) \leq k'_V - \delta$ then $\phi(C \cap D) \leq k_V^0 - (\delta + 1)$ and hence the left valley V_l or the right valley V_r of C, D exists and before the close $\text{Res}(V_l) + \text{Res}(V_r) \geq 2(\delta + 1)$. Since Res is decreased by at most 1, we have after the close $\text{Res}'(V_l) + \text{Res}'(V_r) \geq 2\delta$.

We conclude by showing that Invariant 2 implies Invariant 3. Let V be an arbitrary valley during the realignment of S_{aug} and $A \subseteq V$ a busy interval on processor k_V with $A \sim V$. Let C be the critical set with $C \sim V$. Note that $A \sim V$ implies that A intersects C . Assume for contradiction that A does not span C . The $\min A$ is within subinterval of C or symmetrically $\max A$ within subinterval of C . We assume the first case, i.e. $t := \min A - 1 \in C$ and $\min A \in C$. The second case follows by symmetry. If $t \in V$, then time slot t is busy on processor k_V by Invariant 2. Therefore, A cannot be a (full) busy interval on processor k_V , contradicting the choice of A . If $t = \min V - 1$, then consider valley W with $t \in W$ and $t \sim W$. We must have $W \supset V$, $W \succ V$ and $t \in C_W$. Therefore $k_W \geq k_V$ and Invariant 2 implies that $t = \min A - 1$ is busy on processor k_V , again contradicting the choice of A as full busy interval on processor k_V . \square

Lemma 28. *The result S_{real} of the realignment of S_{aug} is defined.*

Proof. The Invariants in Lemma 27 imply that whenever we are inside the while-loop of fill , the corresponding busy interval A indeed spans the corresponding critical set C with $C \sim A$, hence V_r or V_l exists and there is sufficient Res to close some valley.

Since in the while-loop of $\text{fill}(k, V)$ the busy interval $A \subseteq V$ on k_V always spans C if $V \sim C$ (Invariant 3), the left valley V_l and the right valley V_r of the critical set C and interval A are properly defined. Also since $\phi(A) \leq k - 1$, Invariant 1 implies that V_l or V_r exists and that there is sufficient Res such that one of the two valleys of C is closed in this iteration. This reduces the number of idle intervals on processor k by at least 1, since Invariant 2 implies that V_l or V_r cannot end strictly within an idle interval on k . Hence all terms in the realignment are well defined and the realignment terminates. \square

Lemma 29. *For every processor $k \in [m]$, every busy interval A on processor k in S_{real} with $\text{crit}(A) \geq 2$ we have $\hat{\phi}(A) > k - 1$.*

introduce this notation for critical set C_V of valley V , i.e. $C_V \sim V$.

Proof. We show that $\text{fill}(k, T)$ establishes the property on k . The claim then follows since $\text{fill}(k, T)$ does not change the schedules of processors above k . Since in $\text{fill}(k, T)$, we always close valleys for busy intervals A on k spanning a corresponding critical set C , we know that on processor k , busy intervals are only extended. Let $A \subseteq V$ be a busy interval on processor k in S_{real} with $A \sim V$ and $\text{crit}(A) \geq 2$. No valley $W \supseteq V$ can have been closed on k since otherwise there would be no $A \subseteq V$ in S_{real} . Consider then the point in $\text{fill}(k, V)$ when the while-loop terminates. Clearly at this point all $A' \subseteq V$ with $A' \sim V$ on processor k have $\hat{\phi}(A') > k - 1$. There must also be at least one such A' at this point for $A \subseteq V$ to be a busy interval on k in S_{real} with $A \sim V$. In particular, one such A' must have $A' \subseteq A$, which directly implies $\hat{\phi}(A) \geq \hat{\phi}(A') > k - 1$. \square

Lemma 30. *The realignment of S_{aug} does not create new activation times but may only change the corresponding processor being activated, i.e. if $t \in T$ is an activation of some processor k in S_{real} , then t is also an activation of some processor k' in S_{aug} .*

Proof. Consider the first step in the realignment of S_{aug} in which some $t \in T$ becomes an activation of some processor k' where t was no activation of any processor before this step. This step must be the closing of some valley V on some processor $k > k'$: On processor k , we have seen that closing of some valley can only merge busy intervals. On processors above k , the schedule does not change.

Busy slots on processors $k'' < k$ are only removed (definition close), therefore $t - 1$ must have been busy on processor k' and idle on $k' + 1, \dots, k$ before the close. Refer to Figure 4.3 for a visual sketch of the situation.

Explicitly make this into a separate lemma?

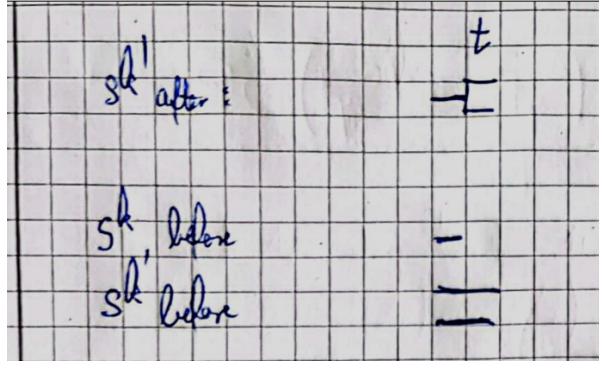


Figure 4.2:

If $t \in V$, then processor $k' + 1$ (or k) must have been busy before at t . Hence t was already an activation before the close, contradicting our initial choice of t .

If $t \notin V$, then $t \succ V$. Let W be the valley s.t. V is closed during $\text{fill}(k, W)$, hence $W \supset V$. If $t \in W$, then $t \sim C_W$ and $t \in C_W$. By Invariant 2, processors $1, \dots, k_W = k$ are busy at t before the close. Again, this implies that t was an activation before the close already, contradicting our choice of t . If $t \notin W$, then let W' be the valley with $t \sim W$ and $t \in W'$. We have $W \prec t \sim W'$ and $W' \supset W$ and $t \in C_{W'}$. Therefore $k_{W'} \geq k_W = k$ and Invariant 2 implies that processors $1, \dots, k$ are busy at t before the close. Hence, t was activation before the close already, again contradicting our initial choice of t . \square

Lemma 31. *Let I an idle interval in S_{real} on some processor k and A_l, A_r the busy intervals to the left and right of I with $\text{crit}(A_l) \leq 1$ and $\text{crit}(A_r) \leq 1$. Allow A_l to be empty, i.e. we might have $\min I = 0$, but A_r must be nonempty, i.e. $\max I < d^*$. Then $k = 1$ and $\hat{\phi}(A_l \cup I \cup A_r) > 0$.*

Proof. By Lemma 30 and $\text{crit}(A_r) \leq 1$, we know that $\min A_r$ is an activation of processor 1 in S_{aug} . Hence $\max I$ is idle in S_{aug} on processor 1 and hence on all processors (stair-property in S_{aug}). Since no jobs are scheduled at $\max I$, we know that $\text{crit}(\max I) \leq 1$ and $J_V = \emptyset$ for all valleys V containing the slot $\max I$, and hence also $\text{Res}(V) = 0$ at all times during the realignment. Therefore, no V intersecting $[\max I, \max A_r]$ was closed during realignment on any processor, since this V would contain $\max I$. Since A_r is a busy interval with $\text{crit}(A_r) \leq 1$ (i.e. not containing activations of processors above 1 in S_{aug}), we must then have $k = 1$.

For I to be idle on processor $k = 1$ in S_{real} and $\text{crit}(I) \geq 2$, some $V \succeq I$ with $V \cap I \neq \emptyset$ and hence $V \supseteq \max I$ would have to been closed, which contradicts what we have just shown. Therefore $\text{crit}(I) \leq q$ and no valley V with $V \cap [\min A_l - 1, \max A_r + 1] \neq \emptyset$ can have been closed. Therefore, the constellation

Properly define what 'the constellation occurs in the same way in S_{aug} ' means.

occurs exactly in the same way in S_{aug} on processor 1, see Figure 4.3.




Figure 4.3:

Let j be the single job scheduled at $\min A_r$. We conclude by showing that $E_j \subseteq I \cup A_r$ and therefore $\hat{\phi}(I \cup A_r) > 0$. Otherwise, j could be scheduled at $\min I$ or $\max A_r + 1$. In the first case, pltr would have extended A_l by scheduling j at time $\min I$ instead of at $\min A_r$. In the second case, pltr would have extended the idle interval I by scheduling j at $\max A_r + 1$ instead of at $\min A_r$. \square

Lemma 32. *For every processor k , every idle interval on processor k in S_{opt} intersects at most two distinct idle intervals of processor k in S_{real} .*

Proof. Let I_{opt} be an idle interval in S_{opt} on processor k intersecting three distinct idle intervals of processor k in S_{real} . Let I be the middle of these idle intervals. Lemma 31 and Lemma 29 imply that k busy processors are required during I and its neighboring busy intervals. This makes it impossible for S_{opt} to be idle on processor k during the whole interval I_{opt} . \square

Sketch this

4.4 Approximation Guarantee

Lemma 33. $\text{costs}(S_{\text{real}}) \leq 2 \text{OPT} + P$

Proof. We first show that $\text{idle}(S_{\text{real}}^k) \leq 2 \text{off}(S_{\text{opt}}^k) + \text{on}(S_{\text{opt}}^k)$ for every processor $k \in [m]$. Let \mathcal{I}_1 be the set of idle intervals on S_{real}^k intersecting some off-interval of S_{opt}^k . Lemma 32 implies that \mathcal{I}_1 contains at most twice as many intervals as there are off-intervals in S_{opt}^k . Therefore, the costs of these idle intervals is bounded by $2 \text{off}(S_{\text{opt}}^k)$.

Let \mathcal{I}_2 be the set of idle intervals on S_{real}^k not intersecting any off-interval in S_{opt}^k . The total length of these intervals is bounded by $\text{on}(S_{\text{opt}}^k)$.

We continue by showing that $\text{busy}(S_{\text{real}}) \leq 2P$. By construction of S_{aug} and the definition of Res and close, we introduce at most as many auxiliary busy slots at every slot $t \in T$ as there are jobs scheduled at t in S_{pltr} . For S_{aug} , an auxiliary busy slot is only added for t with $\text{crit}(t) \geq 2$ and hence $\text{vol}(t) \geq 1$. Furthermore, initially $\text{Res}(V) = 2|J_V|$ for every valley V and $\text{Res}(V)$ is decremented if some V' intersecting V is closed during $\text{fill}(k, T)$. During $\text{fill}(k, T)$ every $t \in T$ is at most closed once (or rather one V' containing t is closed). Finally, auxiliary busy slots introduced by S_{aug} are used in close. In conclusion, we have $\text{costs}(S_{\text{real}}) \leq 2 \text{off}(S_{\text{opt}}) + \text{on}(S_{\text{opt}}) + 2P \leq 2 \text{OPT} + P$. \square

Theorem 34. $\text{costs}(S_{\text{pltr}}) \leq 2 \text{OPT} + P$

Proof. We argue that $\text{costs}(S_{\text{pltr}}) \leq \text{costs}(S_{\text{real}})$. The theorem then follows from Lemma 33. We do this by transforming S_{real} back into S_{pltr} without increasing the costs of the schedule. Removing the auxiliary busy slots clearly cannot increase the costs. Since the realignment of S_{aug} only moves busy slots between processors, but not between different time slots, we can easily restore S_{pltr} (up to permutations of the jobs scheduled on the busy processors at the same time slot) by moving all busy slots back down to the lower numbered processors. By the same argument as in Lemma 2, this does not increase the total costs of the schedule. \square

Define Lemma stair property somewhere

5 Running Time

Theorem 35. *Algorithm pltr has a running time of $O(Fn \log d)$, where F is the time needed for the flow calculation in Figure 2.*

Proof. First observe that every busy interval is created by a pair of calls to `keepIdle` and `keepBusy`. Every call to `keepIdle` searches for the maximal slot $t \in T$ (after the current time slot) such that the maximum-flow calculation in Figure 2 returns a flow of value P . Using binary search this can be done in $O(F \cdot \log(d^*))$ with F denoting the time required for finding the maximum flow.

We bound the number of busy intervals across all processors in S_{pltr} by n . Note that if `keepIdle` returns d^* , then we do not have to calculate `keepBusy` from d^* on. Therefore, the total number of calls to `keepIdle` and `keepBusy` is bounded by $n + m$. If $m > n$ we can restrict our algorithm to use the first n processors only, as there cannot be more than n processors scheduling jobs at the same time. We derive the upper bound of n for the number of busy intervals across all processors by constructing an injective mapping f from the set of busy intervals to the jobs J . For this construction of f we consider the busy intervals in the same order as the algorithm, i.e. from Top-Left to Bottom-Right. We construct f such that $f(A) = j$ for some busy interval A only if $d_j \in A$.

Suppose we have constructed such a partial mapping for busy intervals on processors m, \dots, k up to some busy interval A on k . Let A' be the last plateau with $A' \subseteq A$ and let $k' \geq k$ be the processor of this plateau A' . By construction of f and the choice of A' , there are at most $k' - (k + 1)$ distinct jobs j with $d_j \in [\min A', \max A]$ already mapped to by f . This is since at most $k' - (k + 1)$ busy intervals on processors $k + 1, \dots, m$ intersect the interval $[\min A', \max A]$. Refer to Figure 5 for a visual sketch.

Define plateaus: busy intervals which are completely idle on higher processors




Figure 5.1:

Let C_t be the critical set over activation $t := \min A'$ of processor k' . Let $J' := \{j_1, \dots, j_{k'}\}$ be the k' distinct jobs scheduled at t . We now that $\max A + 1 \notin C_t$ since $\text{vol}(\max A + 1) < k \leq k'$ and $\max A + 1 > t$. Therefore, every job $j \in J'$ with $d_j \geq \max A + 1$ is scheduled at slot $\max A + 1$. Hence there are at least $k' - (k - 1)$ distinct jobs $j \in J'$ with $d_j \in [\min A', \max A]$ and there must be at least one such job j^* which is not mapped to by f so far and which we therefore can use for $f(A)$. \square

References

- A. Antoniadis, N. Garg, G. Kumar, and N. Kumar. Parallel machine scheduling to minimize energy consumption. In *Proceedings of the Thirty-First Annual ACM-SIAM Symposium on Discrete Algorithms*, SODA '20, page 2758–2769, USA, 2020. Society for Industrial and Applied Mathematics.
- R. Graham, E. Lawler, J. Lenstra, and A. Rinnooy Kan. Optimization and approximation in deterministic sequencing and scheduling : a survey. *Annals of Discrete Mathematics*, 5:287–326, 1979. ISSN 0167-5060. doi: 10.1016/S0167-5060(08)70356-X.