
Abstract

We study a particular scheduling setting in which a set of n jobs with individual release times and deadlines has to be scheduled across m homogeneous processors while minimizing the consumed energy. Idle processors can be turned off so as to save energy, while turning them on requires a fixed amount of energy. For the special case of a single processor, the greedy algorithm Left-to-Right guarantees an approximation factor of 2. We generalize this simple greedy policy to the case of multiple processors and show that the energy costs are still bounded by $2\text{OPT} + P$. Our algorithm has a running time of $(n+m) \log(d^*)F$ where d^* is the largest deadline and F the costs of a maximum flow calculation for checking feasibility of an instance.

Contents

1	Algorithm	3
2	Structure of the PLTR-Schedule	5
2.1	Preliminary Definitions	5
2.2	Critical set of time slots	5
2.3	Definitions based on critical sets	7
3	Modification of our Schedule	9
3.1	Augmentation	9
3.2	Realignment	9

1 Algorithm

Algorithm 1 Parallel Left-to-Right

```

 $m_t \leftarrow m$ 
 $l_t \leftarrow 0$ 
for  $k \leftarrow m$  to 1 do
   $t \leftarrow 0$ 
  while  $t < d^*$  do
     $t \leftarrow \text{KEEPIDLE}(k, t)$ 
     $t \leftarrow \text{KEEPACTIVE}(k, t)$ 
function  $\text{KEEPIDLE}(k, t)$ 
  search for maximal  $t' \geq t$  s.t. exists feasible schedule with  $m_{t''} = k - 1 \forall t'' \in [t, t')$ 
   $m_{t''} \leftarrow k - 1 \forall t'' \in [t, t')$ 
  return  $t'$ 
function  $\text{KEEPACTIVE}(k, t)$ 
  search for maximal  $t' \geq t$  s.t. exists feasible schedule with  $l_{t''} = \max\{k, l_{t''}\}k - 1 \forall t'' \in [t, t')$ 
   $m_{t''} \leftarrow k - 1 \forall t'' \in [t, t')$ 
  return  $t'$ 

```

Formal Problem definition, notation

Formulate and format keep-active, keepidle nicer.

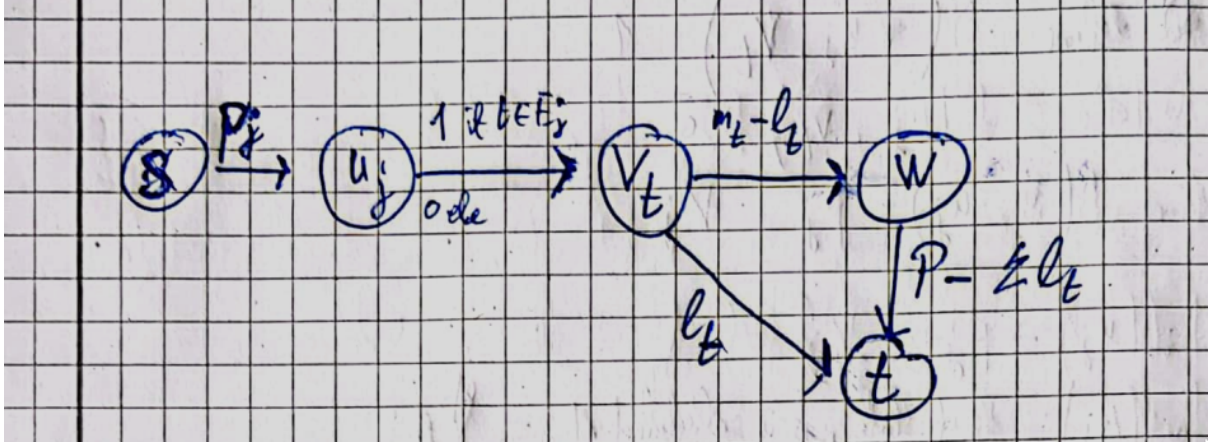
Properly define scheduling problem with lower and upper bounds l_t, m_t .Precisely relate l_t, m_t returned by Algorithm to assignment of jobs to processors, time slots.

Figure 1.1: The Flow-Network for checking feasibility of a scheduling instance with lower and upper bounds l_t and m_t for the number of active processors at t .

Lemma 1. *There exists a feasible solution to a scheduling instance with lower and upper bounds l_t, m_t if and only if the maximum s - t flow in the corresponding flow network depicted in Figure 1 has value P .*

Proof. Let f be a s - t flow of value $|f| = P$. We construct a feasible schedule from f respecting the lower and upper bounds given by l_t and m_t . For every $j \in J$ and $t \in [0, d^*]$, if $f(u_j, v_t) = 1$, then schedule j at slot t . Since $|f| = P$ and the capacity $c(\{s\}, V \setminus \{s\}) = P$, we have $f_{in}(u_j) = p_j$ for every $j \in J$. Hence $f_{out}(u_j) = \sum_{t \in E_j} f_{in}(v_t) = p_j$. Hence every job j is scheduled in p_j distinct time slots.

The schedule respects the upper bounds m_t , since $c(v_t, w) + c(v_t, t) \leq m_t - l_t + l_t$ and for every t at most m_t jobs are scheduled at t .

The schedule respects the lower bounds l_t , since $c(V \setminus \{t\}, \{t\}) = P$ and hence $f(v_t, t) = l_t$ for every slot t . By flow conservation we then have $f_{in}(v_t) \geq l_t$ which implies that at least l_t jobs are scheduled at every slot t .

For the other direction consider a feasible schedule respecting the lower and upper bounds l_t, m_t . We construct a flow f of value P and show that it is maximal.

If j is scheduled at slot t and hence $t \in E_j$, define $f(u_j, v_t) = 1$, otherwise $f(u_j, v_t) = 0$. Define $f(s, u_j) = p_j$ for every $j \in J$. Hence we have $f_{in}(u_j) = p_j$ and $f_{out}(u_j)$ must be p_j since this corresponds to the number of distinct time slots in which j is scheduled. Define $f(v_t, t) = l_t$ for every slot t . Define

Ambiguity of t used for time slots and sink in flow-network: Use α - ω flowBriefly define c for edge capacities and cuts.Introduce notation for number of jobs scheduled at t

$f(v_t, w) = f_{in}(v_t) - l_t$. We have $f(v_t, w) \leq m_t - l_t$ since $f_{in}(v_t)$ corresponds to the number of jobs scheduled at t , which is at most m_t . We also have $f_{out}(v_t) = f_{in}(v_t) - l_t + l_t = f_{in}(v_t)$.

Define $f(w, t) = P - \sum_t l_t$. Then $f_{in}(w) = \sum_t f_{in}(v_t) - l_t = \sum_t |\{j \in J \mid j \text{ scheduled at } t\}| - \sum_t l_t$. Since the schedule is feasible, this corresponds to $f_{in}(w) = P - \sum_t l_t = f_{out}(w)$.

□

2 Structure of the PLTR-Schedule

2.1 Preliminary Definitions

Definition 2. For schedule S , we define the volume $v_S(j, Q)$ of job $j \in J$ in a set $Q \subseteq [0, d^*]$ of time slots as the number of time slots of Q for which j is scheduled at by S .

Definition 3. We define the forced volume $fv(j, Q)$ of job $j \in J$ in a set $Q \subseteq [0, d^*]$ of time slots as the number of time slots of Q for which j has to be scheduled in every feasible schedule, i.e.

$$fv(j, Q) = \max\{0; p_j - |E_j \setminus Q|\}.$$

Definition 4. We define the unnecessary volume $uv_S(j, Q)$ of job $j \in J$ in a set $Q \subseteq [0, d^*]$ of time slots as the amount of volume which does not have to be scheduled during Q , i.e.

$$uv_S(j, Q) = v_S(j, Q) - fv(j, Q).$$

Definition 5. We define the possible volume $pv(j, Q)$ of job $j \in J$ in a set $Q \subseteq [0, d^*]$ of time slots as the maximum amount of volume which j can be feasibly scheduled in Q , i.e.

$$pv(j, Q) = \min\{p_j, |E_j \cap Q|\}.$$

Definition 6. We define the space $space_S(j, Q)$ of job $j \in J$ in a set $Q \subseteq [0, d^*]$ for schedule S as the number of additional time slots, which j can be scheduled in Q , i.e.

$$space_S(j, Q) = pv(j, Q) - v_S(j, Q).$$

Since the corresponding schedule S will always be clear from context, we drop the subscript for v, uv and $space$. We extend our volume definitions to sets $J' \subseteq J$ of jobs by summing over all $j \in J'$, i.e.

$$v(J', Q) = \sum_{j \in J'} v(j, Q).$$

If the first parameter is dropped, we refer to the whole set J , i.e. $v(Q) = v(J, Q)$. Clearly we have for every feasible schedule, every Q, j that $fv(j, Q) \leq v(j, Q) \leq pv(j, Q)$.

Definition 7. We define the density $\phi(Q)$ for a set $Q \subseteq [0, d^*]$ as the average amount of processing volume which has to be completed in every slot of Q , i.e. $\phi(Q) = fv(J, Q)/|Q|$. We also define $\hat{\phi}(Q) = \max_{Q' \subseteq Q} \phi(Q')$.

If $\hat{\phi}(Q) > k - 1$, then clearly at least k processors are required in some time slot $t \in Q$ for every feasible schedule.

Definition 8. We define the deficiency $def(Q)$ of a set $Q \subseteq [0, d^*]$ of time slots as the difference between the amount of volume which has to be completed in Q and the processing capacity available in Q , i.e. $def(Q) = fv(Q) - \sum_{t \in Q} m_t$.

Definition 9. We define the excess $exc(Q)$ of a set $Q \subseteq [0, d^*]$ of time slots as the difference between the processor utilization required in Q and the amount of processing volume available in Q , i.e. $exc(Q) = \sum_{t \in Q} l_t - pv(Q)$.

2.2 Critical set of time slots

Lemma 10. For every s - t cut (S, \bar{S}) we have at least one of the following two lower bounds for the capacity $c(S)$ of the cut: $c(S) \geq P - def(Q(S))$ or $c(S) \geq P - exc(Q(\bar{S}))$, where $Q(S) := \{t \mid v_t \in S\}$.

Proof. Let (S, \bar{S}) be a s - t cut, let $J(S) := \{j \mid u_j \in S\}$. We consider If $w \notin S$, consider the contribution of every node of S to the capacity of the cut.

- Node s : $\sum_{j \in J(\bar{S})} p_j$.
- Node u_j : $|\{v_t \in \bar{S} \mid t \in E_j\}| = |E_j \setminus Q(S)| \geq p_j - fv(j, Q(S))$

- Node v_t : $l_t + m_t - l_t = m_t$

The inequality for node u_j follows since $\text{fv}(j, Q(S)) = \max\{0, p_j - |E_j \setminus Q(S)|\}$. In total, we can lower bound the capacity with

$$\begin{aligned} c(S) &\geq \sum_{j \in J(\bar{S})} p_j + \sum_{j \in J(S)} p_j - \text{fv}(j, Q(S)) + \sum_{t \in Q(S)} m_t \\ &= P - \text{fv}(J(S), Q(S)) + \sum_{t \in Q(S)} m_t \\ &\geq P - \text{def}(Q(S)). \end{aligned}$$

If $w \in S$, again consider the contribution of every node of S to the capacity of the cut.

- Node s : $\sum_{j \in J(\bar{S})} p_j \geq \text{pv}(J(\bar{S}), Q(\bar{S}))$.
- Node u_j : $|E_j \setminus Q(S)| = |E_j \cap Q(\bar{S})| \geq \text{pv}(j, Q(\bar{S}))$
- Node v_t : l_t
- Node w : $P - \sum_t l_t$

In total, we can lower bound the capacity with

$$\begin{aligned} c(S) &\geq P - \sum_{t \in Q(\bar{S})} l_t + \text{pv}(Q(\bar{S})) \\ &= P - \text{exc}(Q(\bar{S})) \end{aligned}$$

□

Lemma 11. *A scheduling instance with lower and upper bounds l_t and m_t is feasible if and only if $\text{def}(Q) \leq 0$ and $\text{exc}(Q) \leq 0$ for every $Q \subseteq [0, d^*]$.*

Proof. If $\text{def}(Q) > 0$ for some Q , then some upper bounds m_t cannot be met. If $\text{exc}(Q) > 0$ for some Q , then some lower bound l_t cannot be met. For the direction from right to left, consider an infeasible scheduling instance with lower and upper bounds. By Lemma 1 we have that the maximum flow f for this instance has value $|f| < P$. Hence, there must be a s - t cut (S, \bar{S}) of capacity $c(S) < P$. Lemma 10 now implies that $\text{def}(Q(S)) > 0$ or $\text{exc}(Q(\bar{S})) > 0$. □

Lemma 12. *For every time slot $t \in [0, d^*]$ for which some processor $k \in [m]$ is activated in S_{pltr} , there exists a set $Q \subseteq [0, d^*]$ of time slots with $t \in Q$,*

$$\begin{aligned} \text{fv}(Q) &= v(Q), \\ v(t') &\geq k - 1 && \text{for } t' \in Q \text{ and} \\ v(t') &\geq k && \text{for } t' \in Q \text{ with } t' \geq t. \end{aligned}$$

Proof. Suppose for contradiction there is some activation $t \in [0, d^*]$ of processor $k \in [m]$ and no such Q exists for t . We show that pltr would have extended the idle interval on processor k which ends at t . Consider the step in pltr when t was the result of keepIdle on processor k and the corresponding lower and upper bounds $m_{t'}, l_{t'}$ for $t' \in [0, d^*]$ right after the calculation of t and the corresponding update of the bounds by keepIdle. We modify the bounds by decreasing m_t by 1. Note that at this point $m_{t'} \geq k$ for every $t' > t$ and $m_{t'} \geq k - 1$ for every t' .

Consider $Q \subseteq [0, d^*]$ s.t. $t \in Q$ and $\text{fv}(Q) < v(Q)$. Before our modification we had $m_Q := \sum_{t' \in Q} m_{t'} \geq v(Q) > \text{fv}(Q)$. The inequality $m_Q \geq v(Q)$ here follows since the upper bounds $m_{t'}$ are monotonically decreasing during pltr. After our modification we still have $m_Q \geq \text{fv}(Q)$.

Consider $Q \subseteq [0, d^*]$ s.t. $t \in Q$ and $v(t') < k - 1$ for some t' . At the step in pltr considered by us, we hence have $m_{t'} \geq k - 1 > v(t')$ and therefore before our decrement of m_t we had $m_Q > v(Q) \geq \text{fv}(Q)$ which implies $m_Q \geq \text{fv}(Q)$ after the decrement of m_t .

Finally, consider $Q \subseteq [0, d^*]$ s.t. $t \in Q$ and $v(t') < k$ for some $t' > t$. Again at the step in pltr considered by us, we have $m_{t'} \geq k > v(t')$ which implies $m_Q \geq \text{fv}(Q)$ after our decrement of m_t .

If for t no Q exists as characterized in the proposition, t cannot have been the result of keepIdle at this step in pltr, which is a contradiction. □

Definition 13. We call such Q for activations t of processor k characterized by Lemma 12 tight set Q_t over activation t of processor k .

Definition 14. A critical set $C_t \subseteq [0, d^*]$ over an activation t is the maximum of the set of tight sets Q_t over activation t in regard to the density ϕ , i.e.

$$C_t := \operatorname{argmax}\{\phi(Q) \mid Q \subseteq [0, d^*] \text{ is tight set over } t\}.$$

As the set of these critical sets C_t for fixed t is closed under union, for the sake of uniqueness, we take C_t to be the inclusion-maximal critical set over activation t .

2.3 Definitions based on critical sets

Definition 15. We define a total order \preceq on the set of critical sets C_t over all activations t . For activations $t, t' \in [0, d^*]$ of processors k and k' respectively, we define $C_t \preceq C_{t'}$ if and only if $k < k'$ or $k = k'$ and $t \geq t'$. In other words, \preceq is the same order in which pltr calculates the activations: from Top-Left to Bottom-Right.

Definition 16. Let $\operatorname{rank} : \{C_t\} \rightarrow \mathbb{N}$ be a mapping to the natural numbers corresponding to \preceq , i.e.

$$\operatorname{rank}(C_t) \leq \operatorname{rank}(C_{t'}) \Leftrightarrow C_t \preceq C_{t'}$$

Definition 17. Let $\operatorname{crit} : \{C_t\} \rightarrow [m]$ be a mapping to the processors s.t.

$$\operatorname{crit}(C_t) = c \Leftrightarrow c \text{ is the highest processor activated at } t$$

Definition 18. We extend these definitions to general time slots $t \in [0, d^*]$.

$$\operatorname{rank}(t) := \begin{cases} \max\{\operatorname{rank}(C) \mid t \in C\} & \text{if } t \in C \text{ for some critical set } C \\ 0 & \text{otherwise} \end{cases}$$

$$\operatorname{crit}(t) := \begin{cases} \max\{\operatorname{crit}(C) \mid t \in C\} & \text{if } t \in C \text{ for some critical set } C \\ 0 & \text{otherwise} \end{cases}$$

We also extend the definitions to intervals $D \subseteq [0, d^*]$.

$$\operatorname{rank}(D) := \max\{\operatorname{rank}(t) \mid t \in D\}$$

$$\operatorname{crit}(D) := \max\{\operatorname{crit}(t) \mid t \in D\}$$

Definition 19. Let C be a critical set. A nonempty interval $V \subseteq [0, d^*]$ is a valley of $\operatorname{rank}(C)$ if $C \sim V$ and V is inclusion maximal. Let C_1, \dots, C_l be the (maximal) intervals of C . A nonempty interval V is a valley of C if V is exactly the interval between C_a and C_{a+1} for some $a < l$, i.e. $V = [\max C_a + 1, \min C_{a+1} - 1]$.

Definition 20. For a critical set C , an interval D spans C if $D \cap C$ contains only full subintervals of C and at least one subinterval of C . The left valley V_l of C and an interval D spanning C is the valley of C ending at $\min(C \cap D) - 1$ (if such a valley of C exists). The right valley V_r of C and an interval D spanning C is the valley of C starting at $\max(C \cap D) + 1$ (if such a valley of C exists).

Definition 21. For a valley V , we define the jobs $J(V) \subseteq J$ as all jobs which are scheduled by S_{pltr} in every $t \in V$.

Lemma 22. For every critical set C with $c := \operatorname{crit}(C)$, every interval D spanning C : if $\phi(C \cap D) \leq c - \delta$ for some $\delta \in \mathbb{N}$, then V_l or V_r is defined and $|J_{V_l}| + |J_{V_r}| \geq \delta$, where we take $|J_V| := 0$ if V does not exist.

Proof. By choice of C as critical set with $c = \operatorname{crit}(C)$ we have $v(C \cap D) \geq (c - 1) \cdot |C \cap D|$. If this inequality is fulfilled strictly, i.e. if $v(C \cap D) > (c - 1) \cdot |C \cap D|$, then with $\operatorname{fv}(C \cap D)/|C \cap D| \leq c - \delta$ we directly get $\operatorname{uv}(C \cap D)/|C \cap D| > \delta - 1$. This implies that there are at least δ jobs j scheduled in $C \cap D$

Provide a rough visual sketch here

Make math more readable in this whole proof, e.g. by using fractions and display math or by replacing division by multiplication

with $uv(j, C \cap D) > 0$. Such jobs must have $E_j \cap (C \setminus D) \neq \emptyset$ and hence at least one of V_l and V_r for C and D must exist and the jobs must be contained in J_{V_l} or J_{V_r} .

If on the other hand we have equality, i.e. $v(C \cap D) = (c - 1) \cdot |C \cap D|$, then let t be the activation of processor c for which C is critical set for. Since $v(t) > c - 1$, we must have $t \notin C \cap D$. By the same argument as before, we have that if $fv(C \cap D)/|C \cap D| \leq c - \delta$, then $uv(C \cap D)/|C \cap D| \geq \delta + 1$. Now suppose that there is no job j scheduled in C s.t. $space(j, C \cap D) > 0$. Then $fv(C \setminus D) = v(C \setminus D) > (c - 1) \cdot |C \cap D|$. Hence $fv(C \setminus D) = v(C \setminus D) > (c - 1) \cdot |C \cap D|$. We then get $\phi(C \setminus D) = v(C \setminus D) > (c - 1) \cdot (C \cap D)$ since by case assumption $t \in C \setminus D$. In conclusion, $C \setminus D$ is still a tight set over t but has higher density than C , contradicting the choice of C . Therefore, there must exist a job j scheduled in C with $space(j, C \cap D) > 0$ and hence

$$\frac{uv(C \cap D) + space(j, C \cap D)}{|C \cap D|} > \delta - 1,$$

which again implies that there must be at least δ jobs scheduled in C with an execution interval intersecting both $C \setminus D$ and $C \cap D$. This implies that the left valley V_l or the right valley V_r of C and D exist and that at least δ jobs are contained in J_{V_l} or J_{V_r} . \square

3 Modification of our Schedule

We modify the schedule S_{pltr} returned by our algorithm in two steps. The first step augments specific processors with auxiliary active slots, s.t. in every critical set C , there are at least the first $\text{crit}(C)$ processors active. Recall that for the single processor ltr algorithm, the crucial property for the approximation guarantee was that every idle interval of S_{OPT} can intersect at most 2 distinct idle intervals of S_{ltr} . The second modification step is more involved and establishes this crucial property on every processor $k \in [m]$ by making use of Lemma 22. It is important to note that these modification steps are only done for the sake of the analysis. By making sure that the costs can only be increased by this modification, we get an upper bound for the costs of S_{pltr} .

Give some high level explanation that we realign the jobs of J_{V_l}, J_{V_r} to higher processors where necessary.

3.1 Augmentation

We transform S_{pltr} into S_{aug} by adding for every t with $k := \text{crit}(t) \geq 2$ and $v(t) = k - 1$ an auxiliary active slot on processor k . This auxiliary active slot does not count towards the volume.

Lemma 23. *In S_{aug} , in every $t \in [0, d^*]$ with $\text{crit}(t) \geq 2$ processors $1, \dots, \text{crit}(t)$ are active.*

Proof. The property directly follows from our choice of the critical sets, the definition of $\text{crit}(t)$ and the construction of S_{aug} . \square

3.2 Realignment

Algorithm 2 Realignment of S_{aug}

```

for  $k \leftarrow m$  to 1 do
  fill( $[0, d^*]$ )
  Res( $V$ )  $\leftarrow$  1 for every  $V$  s.t. some  $V' \cap V \neq \emptyset$  was closed on processor  $k$ 
function fill( $k, V$ )
  if  $\text{crit}(V) \leq 1$  then
    return
  let  $C$  be critical set s.t.  $C \sim V$ 
  while exists active interval  $A \subseteq V$  on processor  $k$  with  $A \sim V$  and  $\hat{\phi}(A) \leq k - 1$  do
    let  $V_l, V_r$  be the left and right valley for  $C$  and interval  $A$  (if  $A$  spans  $C$ )
    if  $V_l$  exists and Res( $V_l$ )  $> 0$  then
      close( $k, V_l$ )
    else if  $V_r$  exists and Res( $V_r$ )  $> 0$  then
      close( $k, V_r$ )
  for every valley  $V' \subseteq V$  of  $C$  which has not been closed on  $k$  do
    fill( $k, V'$ )
function close( $k, V$ )
  for every  $t \in V$  which is idle on processor  $k$  do
    if processors  $1, \dots, k - 1$  idle at  $t$  then
      introduce new auxiliary active slot on processor  $k$  at time  $t$ 
    else
      move active slot at time  $t$  of highest processor among  $1, \dots, k - 1$  to processor  $k$  at  $t$ 

```

References